

Class 8 Mini Project

Marina Puffer (PID: A16341339)

Preparing the data

Data is prepared in CSV format

```
# Save your input data file into your Project directory
fna.data <- "WisconsinCancer.csv"

# Input the data and store as wisc.df. `row.names=1` puts the first column into the row names
wisc.df <- read.csv(fna.data, row.names=1)
# Viewing resulting data frame to check:
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1

	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean
842302	0.11840	0.27760	0.3001	0.14710
842517	0.08474	0.07864	0.0869	0.07017
84300903	0.10960	0.15990	0.1974	0.12790
84348301	0.14250	0.28390	0.2414	0.10520
84358402	0.10030	0.13280	0.1980	0.10430
843786	0.12780	0.17000	0.1578	0.08089

	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419	0.07871	1.0950	0.9053	8.589
842517	0.1812	0.05667	0.5435	0.7339	3.398
84300903	0.2069	0.05999	0.7456	0.7869	4.585
84348301	0.2597	0.09744	0.4956	1.1560	3.445

84358402	0.1809		0.05883	0.7572	0.7813	5.438
843786	0.2087		0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	153.40	0.006399	0.04904	0.05373		0.01587
842517	74.08	0.005225	0.01308	0.01860		0.01340
84300903	94.03	0.006150	0.04006	0.03832		0.02058
84348301	27.23	0.009110	0.07458	0.05661		0.01867
84358402	94.44	0.011490	0.02461	0.05688		0.01885
843786	27.19	0.007510	0.03345	0.03672		0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302	0.03003		0.006193	25.38		17.33
842517	0.01389		0.003532	24.99		23.41
84300903	0.02250		0.004571	23.57		25.53
84348301	0.05963		0.009208	14.91		26.50
84358402	0.01756		0.005115	22.54		16.67
843786	0.02165		0.005082	15.47		23.75
	perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302	184.60	2019.0		0.1622		0.6656
842517	158.80	1956.0		0.1238		0.1866
84300903	152.50	1709.0		0.1444		0.4245
84348301	98.87	567.7		0.2098		0.8663
84358402	152.20	1575.0		0.1374		0.2050
843786	103.40	741.6		0.1791		0.5249
	concavity_worst	concave.points_worst	symmetry_worst			
842302	0.7119		0.2654			0.4601
842517	0.2416		0.1860			0.2750
84300903	0.4504		0.2430			0.3613
84348301	0.6869		0.2575			0.6638
84358402	0.4000		0.1625			0.2364
843786	0.5355		0.1741			0.3985
	fractal_dimension_worst					
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				
84358402		0.07678				
843786		0.12440				

The first column is from a pathologist giving an expert diagnosis, which is essentially the answer to if the cells are malignant or benign, so we should omit it.

```
#Use -1 to remove first column
wisc.data <- wisc.df[,-1]
#Create diagnosis vector to check results later
diagnosis <- as.factor(wisc.df[,1])
```

Exploratory data analysis

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

569 rows, so 569 observations in the dataset.

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
diagnosis
  B    M
357 212
```

There are 357 benign and 212 malignant diagnoses.

Q3. How many variables/features in the data are suffixed with `_mean`?

```
#`colnames()` gives access to the column names
column_names <- colnames(wisc.data)
#`grep()` finds patterns within the names, searches for "_mean" in the column names
column_mean <- grep("_mean", column_names)
#`length()` gives the number of elements in the vector
length(column_mean)
```

```
[1] 10
```

Principal Component Analysis

First check the mean and standard deviation of the features of the data to determine if the data should be scaled.

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data,2,sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se

1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

Execute PCA with `prcomp()` on the data

```
#`prcomp` creates a PCA of the data, and setting "scale" to TRUE ensures that even through
wisc.pr <- prcomp(wisc.data, scale=TRUE)
#there are too many rows to show each individual patient, so use `summary`
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44.27% of the original variance is captured by PC1.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?


PC1-PC3. At PC3, the cumulative proportion is 72.6%.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

PC1-PC7. At PC7, the cumulative proportion is 91%.

Interpreting PCA results

```
biplot(wisc.pr)
```




Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This plot is quite hard to look at, it shows all column names and patient codes, making it extremely difficult to identify points. All the points are jumbled together.

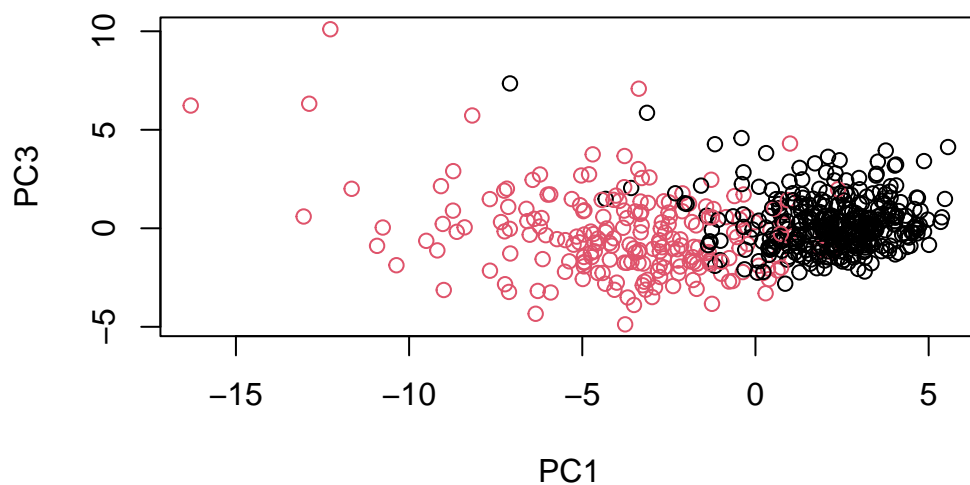
Scatter plot observations by components 1 and 2:

```
plot(wisc.pr$x[,1], wisc.pr$x[,2], col=diagnosis, xlab="PC1", ylab="PC2")
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
plot(wisc.pr$x[,1], wisc.pr$x[,3], col=diagnosis, xlab="PC1", ylab="PC3")
```



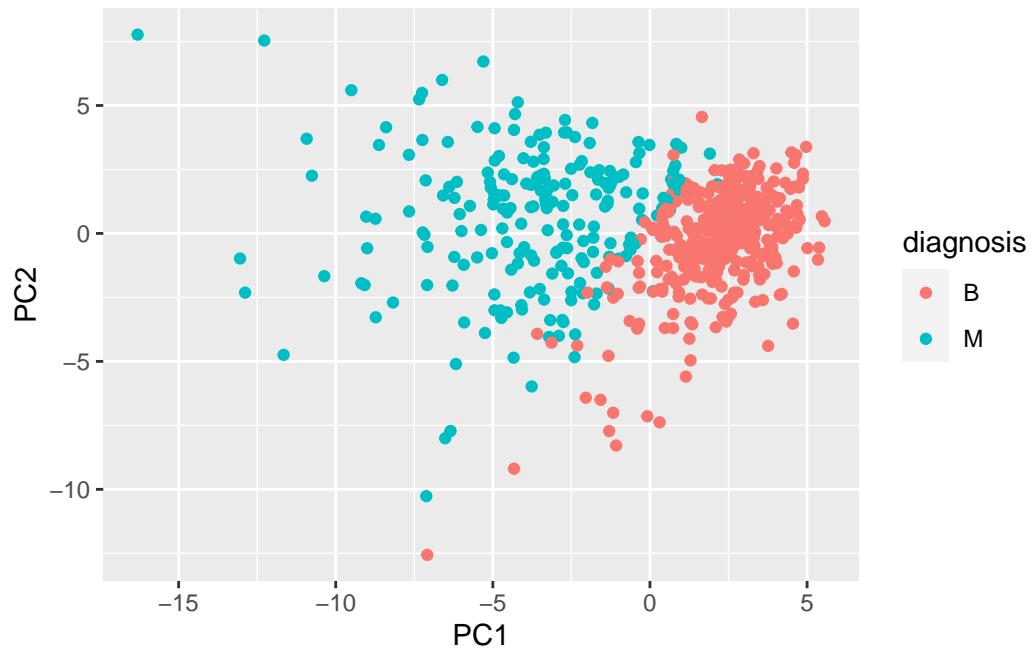
The spread of the points along the y axis is wider than the plot of PC1 vs PC2. There is also more overlap of the benign and malignant points, so the first plot has a cleaner cut separating the subgroups.

Use ggplot2 to make a fancy figure:

```
# Create a data.frame for ggplot
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis

# Load the ggplot2 package
library(ggplot2)

# Make a scatter plot colored by diagnosis
ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```

Variance explained

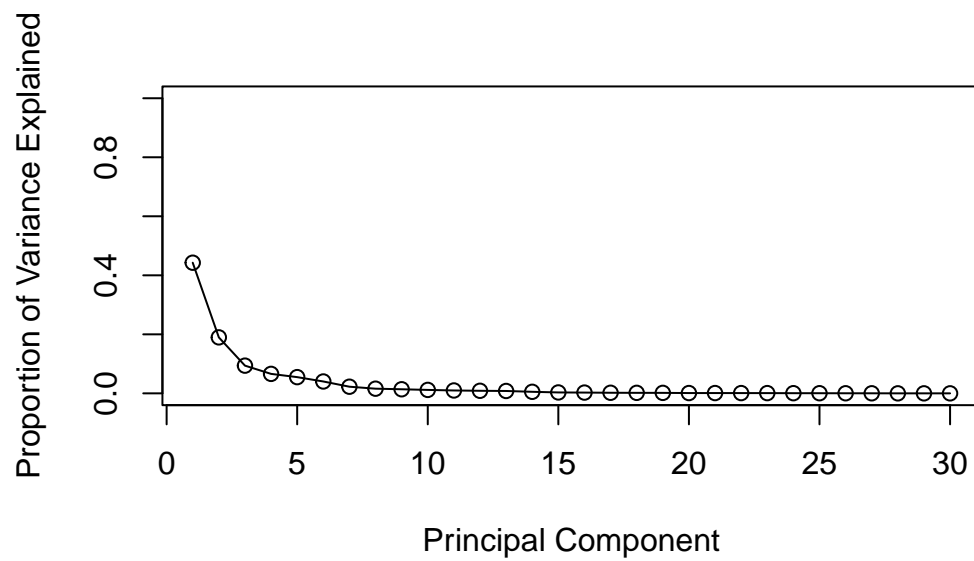
Calculating the variance of each PC by squaring the sdev component of `wisc.pr`

```
pr.var <- (wisc.pr$sdev)^2
head(pr.var)
```

```
[1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```

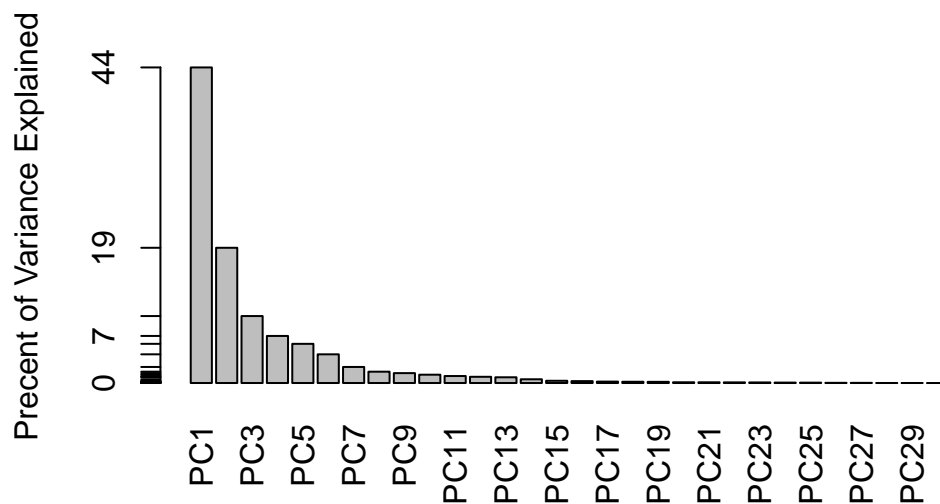
Calculating the variance explained by each PC by dividing by the total variance explained by all PCs.

```
pve <- pr.var/sum(pr.var)
plot(pve, xlab="Principal Component", ylab="Proportion of Variance Explained", ylim=c(0,1))
```



Alternative scree plot of the same data:

```
barplot(pve, ylab = "Precent of Variance Explained",  
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)  
axis(2, at=pve, labels=round(pve,2)*100 )
```



Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
wisc.pr$rotation[,1]["concave.points_mean"]
```

```
concave.points_mean
-0.2608538
```

Component of loading vector for the feature `concave.points_mean` is -0.26

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data?

PC1-PC5.

Hierarchical clustering

This type of clustering does not assume in advance the number of natural groups that exist in the data.

```
#First scale the data
data.scaled <- scale(wisc.data)
#Calculate the distnaces between all pairs of observations
data.dist <- dist(data.scaled)
#Create a hierarchical clustering model using complete linkage
wisc.hclust <- hclust(data.dist, method="complete")
wisc.hclust
```


Call:

```
hclust(d = data.dist, method = "complete")
```

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 569
```

Q11. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```



```
data.dist
hclust (*, "complete")
```

Selecting number of clusters

Use `cutree()` to cut the tree so that it has 4 clusters

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis		
wisc.hclust.clusters	B	M	
1	12	165	
2	2	5	
3	343	40	
4	0	2	

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=5)
table(wisc.hclust.clusters, diagnosis)
```

	diagnosis		
wisc.hclust.clusters	B	M	
1	12	165	
2	0	5	
3	343	40	
4	2	0	
5	0	2	


4 clusters seems to be the best, any increase in the number of clusters only increases the messiness without increasing the number of individuals within each cluster. Decreasing clusters only categorizes both B and M within the same cluster.

Using different methods

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

```
wisc.hclust <- hclust(data.dist, method="complete")
plot(wisc.hclust)
```


Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

```
wisc.hclust.single <- hclust(data.dist, method="single")  
plot(wisc.hclust.single)
```


Cluster Dendrogram



```
data.dist  
hclust (*, "single")
```

```
wisc.hclust.average <- hclust(data.dist, method="average")  
plot(wisc.hclust.average)
```

Cluster Dendrogram



```
wisc.hclust.ward.D2 <- hclust(data.dist, method="ward.D2")  
plot(wisc.hclust.ward.D2)
```

Cluster Dendrogram




```
data.dist  
hclust (*, "ward.D2")
```

The “ward.D2” method gave the best result, it centers the height bars and makes it easy to see that our data can be sectioned into two clusters by the proportions of the height bars.

Combining methods

```
d <- dist(wisc.pr$x[,1:7])  
wisc.pr.hclust <- hclust(d, method="ward.D2")  
plot(wisc.pr.hclust)
```


Cluster Dendrogram




d
hclust (*, "ward.D2")

Generate 2 cluster groups from this hclust object

```
grps <- cutree(wisc.pr.hclust, k=2)
```

```
plot(wisc.pr$x[,1],wisc.pr$x[,2],col=grps)
```



```
table(grps)
```

```
grps
  1  2
216 353
```

```
table(diagnosis)
```

```
diagnosis
  B  M
357 212
```

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

```
table(diagnosis, grps)
```

```
      grps
diagnosis  1  2
  B   28 329
  M  188  24
```

The newly created model separates out the two diagnoses pretty well. Cluster 1 is mostly diagnosed as malignant. Cluster 2 is mostly diagnosed as benign. This can quantify the amount of likely false positives by looking at the individuals within each cluster that do not fall in the majority.